LegacyForward™

LLM ALLIED SKILLS

Participant Guide

For Business Analysts · Quality Analysts · Product Owners
Project Managers · Data Stewards · Change Managers

6-Session Blended Learning Program | v1.0 | March 23, 2026

HOW TO USE THIS GUIDE

This guide is your companion for the full LLM Allied Skills program. It is not a textbook — it is a working document. Write in it, complete the exercises, and return to it during live projects. The annotations you make here will be more valuable than the printed text.

Icon	Means	What to do
📖 **Read**	Core concept — foundational knowledge for all roles	Read carefully. Underline anything that challenges your current practice.
🎯 **Role focus**	Content specific to one or two roles	If this is your role, read closely. Others: skim for awareness.
✏️ **Reflection**	Individual written activity	Write your answer before moving on. Skipping these is the most common learning mistake.
🏗️ **Exercise**	Group or pair activity	Complete with your cohort during the session. Facilitator will debrief.
⚠️ **Watch out**	Common mistake or misconception	This is something your team will get wrong. Discuss it explicitly.
📋 **Job aid**	Template or checklist you keep and use	Copy this into your working files. This is a tool, not a reading.

PROGRAM COMMITMENT

This program works only if you bring real work to it. In every session, you will be asked to apply concepts to a live or recent LLM project at your organization. Come with examples.

© LegacyForward. Participant use only. Not for redistribution.

> Come with failure stories. Those are more useful than success stories.

PROGRAM MAP — 6 SESSIONS

Session	Title	Core Question	Job Aid Produced
1	**Foundational Fluency**	How does an LLM actually work, and why does my current practice break?	*Mental model checklist*
2	**Specification & Scoping**	How do I write requirements when the system's output is probabilistic?	*LLM Acceptance Criteria Template*
3	**Testing & Validation**	How do I test a system I cannot fully predict?	*LLM Test Design Playbook*
4	**Delivery & Iteration**	How do I run a project when velocity is non-linear?	*Meridian Sprint Planner*
5	**Data & Retrieval Quality**	Why does my LLM give different answers to the same question?	*RAG Readiness Checklist*
6	**Change, Governance & Capstone**	How do I govern and communicate a system that can change without a deployment?	*Full Job Aid Library + Credential*

© LegacyForward. Participant use only. Not for redistribution.

SESSION 1 Weeks 1–3 | 3 hours

Foundational Fluency

By the end of this session you will be able to: explain what an LLM is without using technical jargon, and name at least one assumption in your current practice that does not hold for generative AI systems.

1.1 What Is an LLM — Without the Jargon

A Large Language Model is a system trained on enormous amounts of text to predict what comes next in a sequence. It has no memory between conversations, no understanding in the human sense, and no access to information beyond what it was trained on — unless it is connected to a retrieval system.

For practitioners, the two most important things to understand are:

- LLM outputs are probabilistic. Ask the same question twice and you may get different answers. This is not a bug — it is how the system works.
- Quality is a distribution, not a value. An LLM feature does not produce the right answer or the wrong answer. It produces a range of answers with varying quality. Your job is to define what 'good enough' looks like and measure how often the system achieves it.

Token	The unit an LLM processes. Roughly 0.75 words. 'The cat sat' = 4 tokens. Why it matters: every input and output has a token limit. Features that require very long context may hit this limit and silently truncate.
Context window	All the text the model can 'see' at once — your prompt, retrieved documents, and conversation history combined. When this fills up, earlier content is dropped. Why it matters: your requirements need to account for what fits.
RAG	Retrieval-Augmented Generation. The model retrieves relevant documents from a database before generating a response. Why it matters: most enterprise LLM features use RAG. The quality of the retrieval determines the quality of the output — not just the model.
Hallucination	When the model generates confident-sounding text that is factually wrong or fabricated. Why it matters: hallucinations are not always obvious. Testing for them requires different techniques than traditional QA.

© LegacyForward. Participant use only. Not for redistribution.

Temperature	A setting that controls how 'creative' or 'random' the model is. High temperature = more varied outputs. Low temperature = more consistent. Why it matters: production features almost always use low temperature for consistency.

1.2 Why Your Current Practice Breaks

This is the most important section in the entire program. Read it slowly.

THE CORE PROBLEM

Every methodology, process, and mental model you currently use was designed for deterministic systems — systems that behave the same way every time given the same input. LLMs are not deterministic. The gap between what you know and what you need to know is not a skills gap. It is a category gap.

The five assumptions that no longer hold

What you currently assume	Why it breaks with LLMs
Requirements can be written before build begins	You cannot fully specify LLM output behavior before you observe it. Spec and build happen concurrently. Requirements emerge from iteration.
'Done' is a binary state — it passes or it doesn't	LLM output is a distribution. A feature is not done or not done — it hits a confidence band. You need thresholds, not checkboxes.
Velocity is estimable after a few sprints	LLM quality follows an S-curve. Early gains are fast, final gains are expensive. Sprint forecasting based on early velocity will be wrong.
Tests that pass today will pass tomorrow	A model update, corpus change, or prompt edit can silently break previously passing behavior — with no code change. Regression is not additive.
QA validates; the customer accepts at the end	Human judgment must be embedded throughout — not staged at sprint review. The customer's definition of 'good' is an input, not a gate.

⚠️ **WATCH OUT**

The most common reaction at this point is: 'But we can just add some AI-specific steps to our current Agile process.' This is the wrong response. Agile is not wrong — it is not designed for this class of system. Adding steps to a methodology built on broken assumptions does not fix the assumptions.

© LegacyForward. Participant use only. Not for redistribution.

1.3 The Meridian Method™ — What Changes for You

The Meridian Method is the delivery methodology purpose-built for LLM integration in legacy environments. You will learn it in depth across all six sessions. For now, understand one thing: it replaces binary completion with calibration.

Your role	What changes in Meridian	Meridian ceremony you own
Business Analyst	You write behavioral hypotheses, not feature specs. Acceptance criteria become confidence thresholds negotiated against risk tolerance.	*Hypothesis Framing*
Quality Analyst	You characterize output distributions and score against a human baseline. Regression is continuous and non-additive.	*Signal phase + Drift Watch*
Product Owner	You are a human judge — embedded in the build cycle, not reviewing at the end. You define the meridian baseline.	*Meridian Gate*
Project Manager	You plan by calibration cadence, not story points. Progress is measured by confidence delta, not features completed.	*Calibration Standup*
Data Steward	You own retrieval quality. The model's output quality is downstream of your data quality decisions.	*Anchor phase*
Change Manager	You communicate probabilistic behavior to users and own the protocol when the system changes without a deployment.	*Drift Watch communication*

✏️ **Individual Reflection**

Name one thing you currently do in your role that you now realize was built on a deterministic assumption.

What would change about how you write a requirement (or test case, or story, or report) if the system's output could never be fully predicted?

© LegacyForward. Participant use only. Not for redistribution.

What is the biggest practical obstacle you foresee in adopting a calibration-first approach in your current team?

Exercise 1A: The Broken Assumption Hunt

Roles: All roles — mixed groups | Format: Group activity

In groups of 3–4 (mixed roles if possible), discuss a real LLM project you have seen or worked on — or a hypothetical one in your domain.

1. Identify one moment where a broken assumption caused a problem. Describe what happened.
2. Map it to one of the five assumptions from section 1.2.
3. Describe what the team did instead — and whether it worked.
4. Share back: one sentence per group on what the broken assumption cost the project.

© LegacyForward. Participant use only. Not for redistribution.

SESSION 2 Weeks 4–6 | 3 hours

Specification & Scoping

By the end of this session you will be able to: write a behavioral hypothesis for an LLM feature, assign a confidence threshold based on risk class, and explain why a traditional acceptance criterion fails for generative outputs.

2.1 The Problem With 'The System Shall'

Traditional requirements use deterministic language: 'The system shall return the account balance when the user requests it.' This works because the account balance is a fact — it has one correct value.

Now try to write the same kind of requirement for a generative feature: 'The system shall summarize the clinical trial document accurately.' What does accurately mean? Accurate 90% of the time? 99%? How do you measure it? What counts as inaccurate — a wrong fact, a missing fact, a misleading emphasis? Who decides?

These are not questions traditional requirements practice was designed to answer. The Meridian Method replaces the traditional requirement with the Behavioral Hypothesis.

THE BEHAVIORAL HYPOTHESIS

A behavioral hypothesis states: 'We believe the system will [specific behavior] for [defined input class] at [defined confidence threshold], as judged by [specified human calibrators].' Every part of this structure does work. Remove any part and the hypothesis becomes untestable.

Anatomy of a behavioral hypothesis

Component	What it means	Example
Specific behavior	What the system produces — not 'a good summary' but a defined output structure	*Returns a 3-sentence summary that covers: main finding, methodology, and key limitation*
Defined input class	What kinds of inputs this hypothesis covers — edge cases are separate hypotheses	*Clinical trial documents under 20 pages, in English, with a clear abstract*
Confidence threshold	The minimum acceptable success rate, expressed as a percentage	*≥ 85% of outputs rated acceptable by clinical reviewers*

© LegacyForward. Participant use only. Not for redistribution.

Component	What it means	Example
Human calibrators	Who judges quality — named roles, not 'the team'	*Two clinical reviewers from the Medical Affairs team*
Risk class	What happens if the system fails — determines the threshold	*HIGH — outputs inform regulatory submissions; threshold set at 90%*

2.2 Risk Classes and Confidence Thresholds

Not all LLM features carry the same risk. A chatbot that suggests a playlist can fail 20% of the time without consequence. A system that summarizes a legal contract for a human decision-maker cannot. Confidence thresholds must be set against risk class — not against what is technically achievable.

Risk class	Definition	Suggested threshold	Example features
LOW	Failure has negligible consequence; human easily catches error	**≥ 70%**	*Internal search assistants, draft email helpers, idea generators*
MEDIUM	Failure causes rework or user frustration; some downstream impact	**≥ 85%**	*Customer-facing chatbots, document classifiers, meeting summarizers*
HIGH	Failure informs significant decisions; regulatory, financial, or clinical impact	**≥ 92%**	*Contract analysis, clinical note summarizers, compliance checkers*
CRITICAL	Failure has safety, legal, or patient-harm consequences	**≥ 97% + human review mandatory**	*Drug interaction checkers, legal advice tools, safety-critical alerts*

📋 JOB AID 1 — LLM ACCEPTANCE CRITERIA TEMPLATE

Use this template instead of a traditional acceptance criterion for every LLM feature. Complete all fields before development begins. If you cannot complete a field, the feature is not ready to sprint.

Feature name	
Behavioral hypothesis	*We believe the system will ___ for ___ at ___% confidence, as judged by ___*
Input class definition	*What types of input does this hypothesis cover?*

© LegacyForward. Participant use only. Not for redistribution.

Out-of-scope inputs	*What inputs are explicitly excluded from this hypothesis?*
Risk class (LOW / MEDIUM / HIGH / CRITICAL)	
Confidence threshold	*___% (set against risk class, not technical convenience)*
Human calibrators	*Names and roles of the people who will judge quality*
Evaluation criteria	*What specific qualities will judges score? (accuracy, groundedness, relevance, coherence, safety)*
Failure definition	*What does an unacceptable output look like? Describe with 2–3 examples.*
Edge cases (separate hypotheses needed?)	
Hypothesis owner (BA/PO name)	
Review date	

Exercise 2A: Write Your First Behavioral Hypothesis

Roles: BA, PO, QA, PM — all roles | Format: Group activity

Using a real or realistic LLM feature from your domain, complete Job Aid 1 as a group.

5. Choose a feature (BA/PO should lead this choice).
6. Draft the behavioral hypothesis together — every word matters.
7. Assign a risk class and justify it.
8. Name the human calibrators. If you cannot name them, discuss why — and what that tells you about the feature's readiness.
9. Share your hypothesis with another group. Challenge each other: is the threshold too low? Is the input class too broad?

✏️ Individual Reflection

How would you rewrite the last acceptance criterion you produced using the behavioral hypothesis format?

_____

_____

_____

What is the hardest part of writing a confidence threshold for a feature in your domain?

_____

_____

_____

© LegacyForward. Participant use only. Not for redistribution.

© LegacyForward. Participant use only. Not for redistribution.

SESSION 3 Weeks 7–10 | 3 hours

Testing & Validation

By the end of this session you will be able to: design a test approach for an LLM feature that addresses hallucination, behavioral regression, and output quality — without relying on exact-match testing.

3.1 Why Traditional QA Fails

Traditional QA is built on three pillars: you know what the correct output is, the correct output is consistent, and a passing test stays passing. Remove any pillar and the practice collapses. LLM features remove all three.

⚠️ **THE QA TRAP**

The most common QA mistake on LLM projects is writing test cases that look like traditional test cases — 'Input X should produce Output Y.' This approach produces false confidence. The system will pass those tests while silently failing in production on inputs you did not anticipate.

3.2 The LLM Test Design Framework

LLM test design operates across four dimensions simultaneously. Most QA teams only work on one or two. Gaps in any dimension will produce production failures.

Dimension 1: Behavioral Coverage

Does your test set cover the full input class defined in the behavioral hypothesis? This includes typical inputs, edge cases, adversarial inputs (attempts to make the system fail or misbehave), and out-of-scope inputs (to verify the system declines or handles them gracefully).

Dimension 2: Hallucination Testing

Hallucinations are outputs that are fluent and confident but factually wrong. They are not obvious. Testing for them requires human review of a structured sample, comparison against source documents, and specific test cases designed to elicit confabulation — asking the system about things it should not know.

Dimension 3: Behavioral Regression

When the model, corpus, or prompt changes, previously acceptable outputs may degrade without any code change. Behavioral regression testing requires: a fixed set of 'golden' inputs with documented acceptable output ranges, automated scoring on every

© LegacyForward. Participant use only. Not for redistribution.

system change, and a defined confidence delta threshold that triggers a human review.

Dimension 4: Semantic Evaluation

Unlike exact-match testing, semantic evaluation scores the meaning and quality of outputs against defined criteria. Tools: RAGAS (for RAG systems), LLM-as-judge (using a separate model to evaluate), and human scoring rubrics. You need all three — no single approach is sufficient.

3.3 Hallucination Taxonomy

Not all hallucinations are the same. Understanding the type helps you design the right test and route the defect to the right fix.

Type	What it looks like	Root cause & fix
Intrinsic	Output contradicts information in the source documents	Retrieval failure or prompt not instructing the model to stay grounded. Fix: improve retrieval, add grounding instructions.
Extrinsic	Output introduces facts not present in source documents, whether true or false	Model drawing on training knowledge rather than retrieved context. Fix: stricter grounding prompts, citation requirements.
Faithful but wrong	Output is consistent with source but source is wrong	Data quality issue — not a model issue. Fix: data steward involvement, source validation.
Confident refusal	Model says 'I don't have information on that' when it does have relevant retrieved documents	Retrieval quality issue. Fix: improve chunking, embedding, or retrieval ranking.

📋 JOB AID 2 — LLM TEST DESIGN PLAYBOOK

For every LLM feature, complete this checklist before signing off on a test plan. A test plan that cannot answer 'yes' to every item is incomplete.

BEHAVIORAL COVERAGE	
☐	Test set covers ≥ 50 examples from the defined input class

© LegacyForward. Participant use only. Not for redistribution.

☐	Test set includes at least 10 edge cases explicitly documented in the behavioral hypothesis
☐	Test set includes adversarial inputs (prompt injection attempts, out-of-scope requests)
☐	Test set includes out-of-scope inputs to verify graceful handling
☐	Input class boundary cases are documented and tested

HALLUCINATION TESTING	
☐	Structured sample of outputs reviewed against source documents by a human
☐	At least 3 test cases designed specifically to elicit confabulation
☐	Hallucination taxonomy applied: intrinsic, extrinsic, faithful-but-wrong, confident refusal
☐	Defect routing documented: which hallucination type routes to data team vs. prompt team vs. model team

BEHAVIORAL REGRESSION	
☐	'Golden set' of 20–50 fixed inputs documented with acceptable output ranges
☐	Automated scoring runs on golden set after every model update, corpus change, or prompt change
☐	Confidence delta threshold defined: what % drop triggers human review?
☐	Drift Watch alert configured and owner named

SEMANTIC EVALUATION	
☐	Evaluation criteria defined (accuracy, groundedness, coherence, relevance, safety) with scoring rubric
☐	Human scoring rubric validated: two raters agree ≥ 80% of the time on the same outputs
☐	Automated evaluation tool selected (RAGAS / LLM-as-judge / custom) and calibrated against human scores
☐	Evaluation cadence defined: how often, triggered by what, reviewed by whom

Exercise 3A: Test Plan Critique

Roles: QA leads, BA, PO | Format: Group activity

Your facilitator will provide a sample LLM test plan written in traditional QA format.

10. As a group, go through the plan and mark every item that relies on a broken assumption (exact match, binary pass/fail, no regression protocol, etc.).

© LegacyForward. Participant use only. Not for redistribution.

11. Identify which of the four dimensions is missing or underrepresented.
12. Rewrite three test cases from the traditional plan into the Meridian format using Job Aid 2.
13. Discuss: what would this test plan have missed in production?

✏️ **Individual Reflection**

What is the most significant gap between your current test practice and what you just read?

_____

_____

_____

Name a defect from a past project that would have been caught earlier with behavioral regression monitoring.

_____

_____

_____

© LegacyForward. Participant use only. Not for redistribution.

SESSION 4 Weeks 11–14 | 3 hours

Delivery & Iteration

By the end of this session you will be able to: run a Meridian calibration standup, explain why story point velocity does not apply to LLM sprints, and use the Meridian Sprint Planner to structure a 2-week cycle.

4.1 The Calibration Cadence

In Agile, the sprint is a time-boxed delivery unit. In the Meridian Method, the cycle is a calibration unit. The question is not 'what did we complete?' It is 'what did we learn about the system's behavior, and how does it compare to the meridian?'

This requires a different rhythm across every ceremony.

Ceremony	Agile equivalent	What changes in Meridian
Meridian Baseline Session	*Sprint 0*	Human judges document the meridian before any build. Non-negotiable. No baseline = no project start.
Hypothesis Framing	*Sprint Planning*	Feature intent → testable behavioral hypothesis. Every cycle must have a hypothesis or it does not start.
Calibration Standup	*Daily Standup*	'What signal did we observe vs. the meridian?' Human judges present daily — not just at review.
Eval Review	*Sprint Review*	Not a demo. A structured review of eval scores. Output: advance, reset, or retire the hypothesis.
Meridian Gate	*Release Gate*	Human judges confirm quality is real, not a scoring artifact. Binary: gate passed or gate held.
Drift Watch	*No equivalent*	Continuous behavioral regression monitoring. Triggered by any system change, no code change needed.

4.2 Calibration Standup Format

The Calibration Standup replaces the daily standup. The three questions change:

© LegacyForward. Participant use only. Not for redistribution.

❌ Agile standup questions	✅ Meridian standup questions
What did you do yesterday? What will you do today? What is blocking you?	What signal did we observe yesterday? How does it compare to the meridian baseline? What is blocking calibration?

📋 JOB AID 3 — MERIDIAN SPRINT PLANNER

Complete this planner at the start of every Meridian cycle. It replaces the sprint planning board for LLM feature work.

Cycle #	
Dates	*Start: _____ End: _____*
Behavioral hypothesis	*We believe the system will ___ for ___ at ___% confidence, as judged by ___*
Risk class	*LOW / MEDIUM / HIGH / CRITICAL (circle one)*
Current confidence score (start of cycle)	*___%*
Target confidence score (end of cycle)	*___%*
Human judges assigned	*Names:*
Probe approach (what will we build/test?)	
Eval method (how will we score?)	
Definition of advance (what score + gate outcome = next tier?)	
Definition of reset (what would cause us to try a different probe?)	
Definition of retire (what would cause us to abandon this hypothesis?)	
Drift Watch alerts configured?	*YES / NO — if NO, document why*
Data Steward confirmed	*YES / NO — if NO, document dependency*

© LegacyForward. Participant use only. Not for redistribution.

corpus readiness?	

Exercise 4A: Run a Calibration Standup

Roles: PM, BA, PO, QA — all roles | Format: Group activity

In groups of 5–6 (including at least one person in each of: PM, BA/PO, QA, and one 'human judge' role), simulate a Calibration Standup.

Use this scenario: You are on day 3 of a Probe cycle for a contract clause extraction feature. The eval scores from yesterday showed 74% accuracy against the meridian baseline. Target is 85%. You identified 3 inputs where the system missed key obligations.

14. PM facilitates using the three Meridian standup questions.
15. BA/PO: report on signal observed vs. meridian.
16. QA: characterize the distribution — where is it failing?
17. Human judge: react to the 3 missed inputs — are these acceptable misses or blocking failures?
18. PM: close the standup with a clear next action.

Debrief: what felt different from a regular standup? What was harder? What was more useful?

✏️ **Individual Reflection**

What would you change about your last sprint planning session if you had used the Meridian Sprint Planner?

Who in your current team would be the most resistant to the Calibration Standup format? How would you address that?

© LegacyForward. Participant use only. Not for redistribution.

SESSION 5 Weeks 15–18 | 3 hours

Data & Retrieval Quality

By the end of this session you will be able to: explain how RAG pipeline quality affects LLM output, identify the three most common retrieval failures, and complete a RAG Readiness Assessment for a real corpus.

5.1 Why the Model Is Often Not the Problem

When an LLM feature produces poor outputs, the instinct is to blame the model. In most enterprise legacy environments, the model is not the problem. The retrieval pipeline is.

THE DATA STEWARD'S LEVERAGE POINT

In a RAG system, the model can only work with what the retrieval pipeline gives it. If the retrieved documents are irrelevant, incomplete, or misleading, the model will produce poor outputs regardless of how good it is. Data quality is upstream of model quality. This is the Data Steward's leverage point.

5.2 How RAG Actually Works

Retrieval-Augmented Generation has three stages. Each stage is a failure point.

Stage 1: Chunking — *Owner: Data Steward*

Your source documents are split into smaller pieces called chunks. Chunk size and strategy determine what the model can 'see.' A chunk that splits a sentence across two pieces loses meaning. A chunk that is too long may include irrelevant content that confuses the model. Chunking strategy is a data quality decision — not a model decision.

Stage 2: Embedding & Indexing — *Owner: Data Steward + Engineering*

Each chunk is converted into a mathematical representation (an embedding) and stored in a vector database. When a user asks a question, the question is also embedded and the most similar chunks are retrieved. The quality of the embeddings determines the quality of the match. If your documents contain jargon, abbreviations, or domain-specific language, standard embedding models may not represent them well.

Stage 3: Retrieval & Generation — *Owner: QA + Engineering*

The top N chunks are passed to the model along with the user's question. The model

© LegacyForward. Participant use only. Not for redistribution.

generates a response based on what it was given. If the retrieved chunks are wrong, irrelevant, or contradictory, the model will generate a poor response — sometimes confidently. This is why 'faithful but wrong' hallucinations are so hard to catch.

5.3 The Three Most Common Retrieval Failures

Retrieval Failure 1: Semantic Mismatch
The user's question and the relevant document use different words for the same concept. The embedding model returns documents that are lexically similar but semantically irrelevant. Example: a user asks about 'adverse events' but the corpus uses 'safety signals.' Fix: domain-adapted embeddings, synonym expansion, or hybrid search (semantic + keyword).

Retrieval Failure 2: Chunk Boundary Loss
A key piece of information is split across two chunks. The retrieval system returns one chunk but not the other. The model generates a response that is missing critical context. Fix: overlap-aware chunking, larger chunk sizes for dense technical documents, or hierarchical chunking.

Retrieval Failure 3: Stale Index
The source documents have been updated but the vector index has not been refreshed. The model retrieves outdated information and presents it as current. This is particularly dangerous in regulated environments where document versions matter. Fix: index refresh triggers on document update, version-aware chunking, date metadata filtering.

📋 JOB AID 4 — RAG READINESS ASSESSMENT

Complete this assessment before any LLM feature enters the Probe phase. If you cannot answer YES to all items in sections 1 and 2, the corpus is not ready to sprint against.

SECTION 1 — SOURCE DOCUMENT QUALITY	
☐	Source documents have been reviewed for accuracy and are the authoritative versions
☐	Document versions are tracked — it is clear which version was indexed
☐	Domain-specific terminology, acronyms, and jargon are documented in a glossary

© LegacyForward. Participant use only. Not for redistribution.

☐	Documents are in a consistent format (not scanned PDFs, images, or poorly structured HTML)
☐	Sensitive or restricted documents are identified and access controls are applied to the index

SECTION 2 — CHUNKING & INDEXING	
☐	Chunking strategy is documented (size, overlap, method)
☐	Chunk boundaries have been manually reviewed on a sample of 20+ documents
☐	No key facts are split across chunk boundaries in the reviewed sample
☐	Embedding model is appropriate for the domain (tested against domain-specific queries)
☐	Index refresh protocol is documented — who triggers it, when, and how it is verified

SECTION 3 — RETRIEVAL QUALITY	
☐	Retrieval tested against 20+ known queries with verified correct documents
☐	Retrieval precision measured: what % of returned chunks are relevant?
☐	Retrieval recall measured: what % of relevant chunks are returned?
☐	Failure cases documented: what queries consistently retrieve wrong documents?
☐	Hybrid search considered for domain-jargon-heavy corpora

SECTION 4 — GOVERNANCE	
☐	Data Steward named and accountable for corpus quality
☐	Process defined for adding, updating, or removing documents from the corpus
☐	Drift Watch configured to alert on index changes
☐	Audit trail exists for who changed what in the corpus and when

✏️ **Individual Reflection**

Think of a data source your organization would want to use in an LLM feature. Which retrieval failure is most likely? Why?

_____

_____

_____

Who in your organization would own the RAG Readiness Assessment? Is that person currently involved in LLM projects?

_____

_____

_____

© LegacyForward. Participant use only. Not for redistribution.

© LegacyForward. Participant use only. Not for redistribution.

SESSION 6 Weeks 19–24 | 4 hours

Change, Governance & Capstone

By the end of this session you will be able to: communicate probabilistic system behavior to a non-technical stakeholder, apply all four job aids to a real or simulated project, and demonstrate Meridian ceremony participation in your role.

6.1 Communicating Probabilistic Behavior

One of the hardest change management challenges in LLM adoption is explaining to users and executives why the system 'changed' — when in fact there was no deployment, no code change, and no visible event. A model update or corpus refresh is invisible. Its effects are not.

THE CHANGE MANAGER'S CHALLENGE

Users expect software to be consistent. LLM systems are not consistently consistent. Your job is to build a mental model in your organization that separates 'this is broken' from 'this is behaving within its probabilistic range' — and to communicate when the range has shifted.

6.2 AI Incident Taxonomy

Not all LLM failures are the same kind of failure. Before escalating, categorize the incident correctly — the routing determines the fix.

Incident type	Description	Route to
Model drift	Output quality has degraded after a model version update with no other changes	Engineering + QA: re-run Drift Watch against golden set, determine if threshold breach requires rollback
Corpus drift	Output quality has degraded after documents were added, changed, or removed from the retrieval corpus	Data Steward: review recent corpus changes, re-run RAG Readiness Assessment sections 2–3
Prompt regression	A prompt change introduced by the team has degraded output quality	Engineering + QA: compare eval scores before and after change, revert if breach
Distribution shift	User inputs have changed character — new use cases, languages, or domains the system was not calibrated for	BA/PO + QA: assess whether the new input class is within scope, may require new behavioral hypothesis
Within-range variation	User reports a bad output but investigation shows it is within the defined confidence band	Change Manager + PO: communicate expected behavior, review whether threshold is appropriate for user

© LegacyForward. Participant use only. Not for redistribution.

Incident type	Description	Route to
		expectations
True defect	Systemic failure outside confidence band with identifiable root cause	Standard defect process + Drift Watch alert + Meridian Gate re-run before re-release

6.3 Governance Accountabilities

The Meridian Method distributes governance responsibility across roles. This is deliberate — governance of a probabilistic system cannot sit in one place.

Role	Governance accountability
Business Analyst	Owns the behavioral hypothesis register. Ensures every feature in production has a documented, current hypothesis with a named risk class and threshold.
Quality Analyst	Owns the golden test set and Drift Watch configuration. Reports confidence scores at every Eval Review. Escalates threshold breaches.
Product Owner	Acts as human judge at Meridian Gates. Signs off on confidence thresholds before production release. Cannot delegate this to engineering.
Project Manager	Owns the calibration cadence. Ensures Meridian ceremonies run on schedule. Reports on confidence delta progress, not feature velocity.
Data Steward	Owns corpus quality and the RAG Readiness Assessment. Signs off before every Probe cycle that depends on retrieval. Owns index refresh governance.
Change Manager	Owns user communication about system behavior and behavioral changes. Maintains the AI Incident taxonomy log. Runs post-incident user communication.

CAPSTONE — FULL MERIDIAN CYCLE SIMULATION

The capstone runs across the final two weeks of the program. Each cohort group applies everything learned to a real or simulated LLM project. Every role participates in every ceremony. The facilitator observes and scores.

Exercise 6A: Full Meridian Cycle — Capstone
Roles: All roles — full cohort group | Format: Group activity

Your group will simulate one complete Meridian cycle for an LLM feature in your domain. All

© LegacyForward. Participant use only. Not for redistribution.

six ceremonies must be run. All four job aids must be completed.

Ceremony	What to produce
Meridian Baseline Session (DS + PO)	Documented human meridian: 5 examples of acceptable output, 5 examples of unacceptable output, with written rationale for each.
Hypothesis Framing (BA)	Completed Job Aid 1 — behavioral hypothesis, risk class, threshold, calibrators.
Calibration Standup (PM)	Run a 10-minute Calibration Standup using the three Meridian questions. Record the output.
Eval Review (QA)	Using simulated eval scores provided by facilitator, run the Eval Review. Outcome: advance, reset, or retire — with rationale.
Meridian Gate (PO + human judges)	Run the Gate. Human judges review a sample. Binary outcome: gate passed or gate held.
Drift Watch (QA + CM)	Define the Drift Watch configuration: what triggers an alert, who receives it, what is the response protocol.

PROGRAM COMPLETION & CREDENTIAL

Participants who complete all six sessions, all four job aids, and the capstone exercise are eligible for the LLM Allied Skills Certified Practitioner credential issued by LegacyForward.

Credential	LLM Allied Skills Certified Practitioner™
Issued by	LegacyForward
Role tracks	BA Track / QA Track / PO Track / PM Track / Data Steward Track / Change Manager Track
Renewal	Annual — requires evidence of active Meridian Method participation
Assessment	Meridian Method™ Practitioner Assessment — role-specific scenario questions + job aid submission
Recognition	Included in LegacyForward Certified Practitioner Registry

LLM Allied Skills | Participant Guide

Powered by the Meridian Method™

legacyforward.ai | © 2026 LegacyForward. Participant use only.

© LegacyForward. Participant use only. Not for redistribution.