LegacyForward

THE

MERIDIAN METHOD™

A Calibration-First Development Methodology

for LLM Integration in Legacy Enterprise Environments

White Paper | v1.0 | 2025 | legacyforward.ai

EXECUTIVE SUMMARY

Agile did not replace Waterfall. It accelerated it. Both methodologies share a foundational assumption that makes them structurally incompatible with LLM-based system development: the belief that requirements can be fully specified before the system is built, that 'done' is a binary state, and that a passing test today will still pass tomorrow.

These assumptions are false for generative AI systems. LLM outputs are probabilistic distributions, not deterministic values. Behavior that passes at 80% confidence today may degrade to 60% after a model update, a RAG corpus change, or a shift in the input distribution — with no code change and no visible signal.

The Meridian Method™ is a calibration-first development methodology designed specifically for LLM integration in legacy enterprise environments. It replaces binary completion gates with confidence thresholds, replaces sprint velocity with calibration cadence, and elevates human judgment from a final reviewer to a first-class instrument throughout the delivery cycle.

THE CORE CLAIM

The Meridian Method™ is not a patch on Agile. It is a purpose-built methodology for a class of systems that did not exist when Agile was conceived. Its organizing principle: human judgment is the meridian — the fixed reference line from which all system behavior is measured, toward which every iteration converges.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

SECTION 1 — WHY AGILE BREAKS ON LLM PROJECTS

Agile was designed for deterministic systems. Every ceremony, artifact, and principle in Agile assumes that if you build the right thing correctly, it will behave consistently. That assumption has never been fully true — but it was close enough to be useful for traditional software. For LLM-based systems, it is not close enough.

The following five foundational Agile assumptions fail in LLM delivery environments. Each failure is not a matter of degree — it is a categorical breakdown that requires a different methodology, not a modified process.

Assumption 1: Requirements Can Be Specified Before Build

Agile relaxed Waterfall's big-bang specification, but it retained the assumption that a story can be written, pointed, built, and accepted in a bounded cycle. The BA writes the requirement; the team builds to it; QA validates against it.

In LLM delivery, this sequence inverts. You cannot fully specify output behavior before you observe it. A BA writing acceptance criteria for a generative summarization feature cannot know in advance how the model will handle edge cases, ambiguous inputs, or low-quality retrieval results — until the system is running. Specification and build are not sequential phases. They are concurrent, iterative, and mutually dependent.

> **ASSUMPTION FAILURE**
> In LLM delivery, specification is not an input to the sprint. It is an output of it. Requirements emerge from observed system behavior — they cannot fully precede it.

Assumption 2: 'Done' Is a Binary State

Agile's Definition of Done is one of its most powerful constructs — a shared, explicit agreement that a feature is complete and acceptable. It works because traditional systems either behave correctly or they do not. The acceptance criterion is satisfied or it is not.

LLM output is not binary. It is a distribution. A feature that returns a correct, well-formed, contextually appropriate response 87% of the time is not 'done' or 'not done' — it occupies a confidence band. Whether that band is acceptable depends on the use case, the risk tolerance, and the consequences of the 13% failure rate. No existing Definition of Done framework accommodates this. Teams using Agile on LLM projects either accept features at artificially low quality thresholds or never reach 'done' at all.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

ASSUMPTION FAILURE

Done is not a binary state for generative systems. It is a confidence threshold negotiated against risk tolerance. The Meridian Method replaces Definition of Done with Confidence Gate — a defined minimum threshold that varies by feature risk class.

Assumption 3: Velocity Is Estimable

Sprint velocity in Agile assumes that once a team understands the work, effort is roughly estimable and stabilizes over time. Story points, capacity planning, and release forecasting all depend on this predictability.

LLM development has non-linear iteration curves. A prompt pipeline that performs at 75% accuracy may require two sprints to reach 85% and six more to reach 92% — or it may never reach 92% with the current architecture and require a fundamental redesign. There is no way to estimate this upfront. Eval-driven improvement does not follow a linear effort curve. Teams that apply Agile velocity planning to LLM projects will chronically misforecast timelines and exhaust stakeholder confidence before reaching production-grade behavior.

ASSUMPTION FAILURE

LLM iteration follows an S-curve, not a linear effort curve. Early gains are fast; the final 10–15% of quality improvement consumes disproportionate effort. Velocity-based planning cannot accommodate this. The Meridian Method replaces velocity with calibration cadence — progress is measured by confidence delta, not story points completed.

Assumption 4: Regression Is Additive

In Agile, a passing test stays passing unless code changes. Regression suites are additive — once green, always green unless a developer touches the relevant code. This is so fundamental to how QA is practiced that most teams never examine the assumption.

In LLM systems, regression is not additive. A model update, a change to the RAG corpus, a shift in the embedding model, or even a change in the input distribution can alter output behavior across features that were previously stable — with no code change, no deployment, and no visible signal. A QA team running traditional regression on an LLM system will sign off on releases that have silently degraded. The defects will surface in production, attributed to the model rather than to the methodology that failed to catch them.

ASSUMPTION FAILURE

LLM systems can regress without any code change. Traditional regression testing cannot detect this. The Meridian Method introduces Behavioral Regression Monitoring — continuous eval scoring against fixed behavioral baselines, triggered by any system

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

change including model updates and corpus changes.

Assumption 5: The Team Builds; the Customer Accepts

Agile introduced closer customer collaboration than Waterfall, but the boundary remains: the team builds the system, the product owner accepts or rejects it. Human judgment enters at the end of the cycle — in sprint review, in acceptance testing, in production feedback loops.

LLM systems require human judgment to be woven into the build process continuously, not staged at the boundary. The quality of a generative output cannot be evaluated by a test suite alone — it requires human calibration of what 'good' means in context. If that calibration is deferred to sprint review, the team has spent two weeks building toward the wrong definition of quality. Human judgment is not a final reviewer in LLM delivery. It is the instrument that defines the target.

ASSUMPTION FAILURE

Human judgment is not a gate at the end of the LLM delivery cycle. It is the meridian — the reference instrument that must be established first, consulted continuously, and treated as a first-class input throughout. This is the central organizing principle of the Meridian Method.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

SECTION 2 — WATERFALL vs. AGILE vs. MERIDIAN METHOD™

The following table maps the key delivery dimensions across all three methodologies. The Meridian Method is not positioned as superior to Agile for all software development — it is purpose-built for the specific challenges of LLM integration in legacy environments.

Dimension	Waterfall	Agile	Meridian Method™
Unit of work	Requirement document	User story	**Behavioral hypothesis**
Completion signal	Signed-off deliverable	Definition of Done	**Confidence Gate (threshold met)**
Progress measure	Milestone completion	Velocity (story points)	**Calibration delta (confidence change per cycle)**
Requirements timing	Before build begins	Sprint start (just-in-time)	**Concurrent with build; emerge from observation**
Human judgment role	Sign-off at phase gates	Sprint review & acceptance	**First-class instrument; continuous throughout**
Regression model	Additive test suite	Additive automated tests	**Behavioral regression monitoring; non-additive**
Failure mode	Late discovery; costly rework	Incomplete features; velocity misforecast	**Silent degradation; eval drift; misdiagnosed defects**
Legacy environment fit	Poor (rigid, slow)	Partial (better iteration, same spec assumptions)	**Designed for legacy: noisy data, tribal knowledge, fragmented integration**
Team structure	Siloed by phase	Cross-functional squad	**Calibration-integrated: human judges embedded in every cycle**
Model update handling	Not applicable	Not applicable	**Mandatory re-calibration trigger on any system change**
Primary IP home	PMI / PMBOK	Scrum Alliance / SAFe	**LegacyForward™**

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

SECTION 3 — THE MERIDIAN CEREMONIES

The Meridian Method defines six ceremonies that replace or substantially modify their Agile equivalents. Each ceremony is designed around the calibration-first principle: human judgment is established and consulted before, during, and after every build cycle — not staged at the end.

Ceremony	Replaces	Who Attends	Cadence	Output / Purpose
Meridian Baseline Session	*Sprint 0 / Project Kickoff*	BA, PO, QA, Domain SMEs, human judges	Once per initiative	Establishes the Human Meridian: documented examples of acceptable vs. unacceptable outputs across all feature risk classes. This is the fixed reference point for all subsequent calibration.
Hypothesis Framing	*Sprint Planning*	BA, PO, Tech Lead, QA	Start of each cycle (1–2 weeks)	Converts feature intent into a testable behavioral hypothesis: 'We believe the system will [behavior] for [input class] at [confidence threshold].' No hypothesis = no cycle start.
Calibration Standup	*Daily Standup*	Full team including human judges	Daily	Replaces 'what did you do / what will you do' with 'what signal did we observe / how does it compare to the meridian / what is blocking calibration.' Keeps human judgment in the daily loop.
Eval Review	*Sprint Review*	Full team + stakeholders	End of each cycle	Reviews eval scores against confidence thresholds. Not a demo of features — a structured review of behavioral evidence. Outputs: advance, reset, or retire the hypothesis.
Meridian Gate	*Sprint Retrospective + Release Gate*	PO, QA, human judges, risk owner	At confidence threshold or release decision point	The formal human calibration check before any feature advances to the next confidence tier or moves to production. Binary outcome: gate passed or gate held.
Drift Watch	*No Agile equivalent*	QA, MLOps, Data Steward	Continuous / event-triggered	Ongoing behavioral regression monitoring. Triggered automatically by any system change: model update, corpus change, embedding model change, prompt modification. Alerts if confidence delta exceeds defined threshold.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

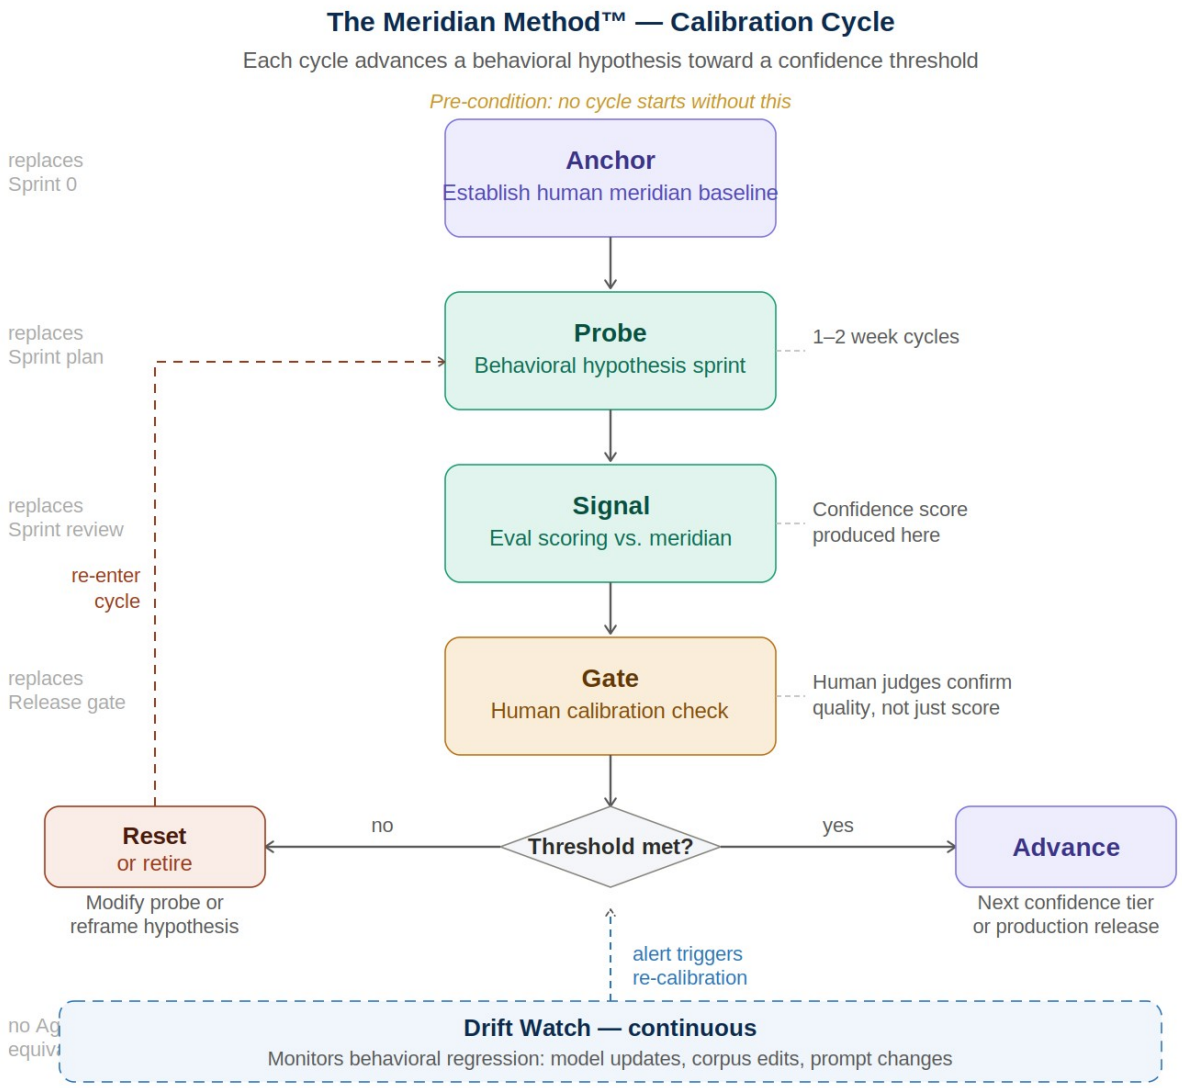

SECTION 4 — THE MERIDIAN SPRINT MODEL

The Meridian sprint is not a time-boxed delivery unit. It is a calibration cycle — a structured loop that moves a behavioral hypothesis from statement to evidence to decision. Each cycle has five phases. Progress is not measured by story points completed but by the change in confidence score relative to the meridian baseline.

ANCHOR	PROBE	SIGNAL	GATE	ADVANCE or RESET
Establish human meridian baseline	*Behavioral hypothesis sprint (1–2 weeks)*	*Eval review: confidence scored*	*Meridian Gate: human calibration*	*Progress or re-hypothesis*

← Each cycle repeats until Confidence Threshold is met or hypothesis is retired →

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

Phase 1 — ANCHOR

Before any build begins, the team establishes or confirms the human meridian for the feature under development. This means reviewing the documented examples from the Meridian Baseline Session and confirming that the human judges for this cycle are available, calibrated, and aligned on what 'good' means for this specific behavioral hypothesis.

If the meridian has not been established for this feature class, the cycle does not start. Anchor is a non-negotiable gate. This is the most significant departure from Agile — it is structurally impossible to begin a Meridian cycle without first establishing the human reference point.

Phase 2 — PROBE

The team builds the minimum viable prompt pipeline, retrieval configuration, or model integration needed to generate observable output for the hypothesis. The goal is not a complete feature — it is a signal. Build only enough to produce evidence. Probe cycles are short: one to two weeks maximum.

This phase is where engineers, prompt engineers, and data stewards are most active. BA and QA roles are present but observational — their job in this phase is to document observed behavior against the hypothesis, not to direct the build.

Phase 3 — SIGNAL

The team runs the eval framework against the probe outputs. Every output is scored against the meridian baseline using defined evaluation criteria: accuracy, groundedness, coherence, relevance, and risk-class-specific criteria. The output of this phase is a confidence score and a behavioral distribution — not a pass/fail.

The Signal phase is where the QA role is most differentiated from traditional testing. QA in Meridian is not looking for defects in the traditional sense — they are characterizing a distribution and mapping it against the meridian. This requires LLM eval fluency, not traditional test automation skills.

Phase 4 — GATE

The Meridian Gate is the human calibration check. Human judges — domain experts, product owners, or designated calibrators — review a structured sample of outputs and their eval scores. They confirm, challenge, or adjust the scoring. They decide whether the confidence score represents genuine quality or a scoring artifact.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

The Gate is not a demo. It is not a sprint review. It is a structured instrument check — the team is verifying that the meridian itself is still accurate, not just that the system is hitting a number. This distinction is critical: a system can hit a confidence threshold while the threshold is wrong.

Phase 5 — ADVANCE or RESET

If the Gate confirms that the confidence threshold has been met for this cycle, the hypothesis advances: the team moves to the next confidence tier, the next feature, or — if the final threshold is met — to the Meridian Gate for production release.

If the Gate does not confirm, the team makes one of three decisions: Reset the probe with a modified approach; Reframe the hypothesis if the behavioral target was wrong; or Retire the hypothesis if the evidence suggests the feature is not achievable at acceptable confidence within the current system architecture. Retirement is not failure — it is honest signal. Agile has no equivalent construct.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

SECTION 5 — MERIDIAN AND THE ALLIED SKILLS PROGRAM

The Meridian Method™ defines what practitioners need to do differently. The LLM Allied Skills Program defines what they need to know to do it. The two are intentionally designed as a system — neither is sufficient without the other.

Role	Meridian Ceremony They Own	Allied Skills Training Module
Business Analyst	Hypothesis Framing — converts feature intent into testable behavioral hypotheses	*Phase 2: Specification & Scoping — writing acceptance criteria for probabilistic outputs*
Quality Analyst	Signal phase execution + Drift Watch monitoring	*Phase 3: Testing & Validation — LLM eval frameworks, hallucination taxonomy, behavioral regression*
Product Owner	Meridian Gate — confirms confidence thresholds and gate decisions	*Phase 2: Specification & Scoping — understanding confidence bands, risk classes, DoD reframe*
Program / Project Manager	Calibration Standup facilitation + cycle planning	*Phase 4: Delivery & Iteration — eval-driven sprint model, calibration cadence, progress measurement*
Data Steward	Anchor phase — confirms RAG corpus quality before each cycle	*Phase 5: Data & Retrieval Quality — chunking strategy, retrieval relevance, corpus governance*
Change Manager	Drift Watch communication — translating behavioral changes to user impact	*Phase 6: Change, Governance & Capstone — AI incident taxonomy, communicating probabilistic behavior*

The implication is direct: an organization can adopt the Meridian Method without the Allied Skills training program, but the ceremonies will be performed incorrectly by practitioners who lack the underlying fluency. Conversely, practitioners who complete the Allied Skills program but operate inside Agile will find their new skills have no structural home. The methodology and the training program are designed to be sold and adopted together.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

SECTION 6 — INTELLECTUAL PROPERTY & LICENSING

IP NOTICE
The Meridian Method™ is proprietary intellectual property of LegacyForward. All rights reserved. This document is a confidential draft. Unauthorized reproduction, distribution, or adaptation of any methodology component, ceremony definition, or framework element is prohibited.

6.1 What Is Owned

- The Meridian Method™ name and trademark
- The five-phase sprint model (Anchor → Probe → Signal → Gate → Advance/Reset)
- The six ceremony definitions and their structured formats
- The Confidence Gate construct and its risk-class tiering framework
- The Behavioral Regression Monitoring protocol and trigger taxonomy
- The Human Meridian Baseline methodology and documentation format
- The Hypothesis Framing template and evaluation rubric
- All associated job aids, templates, and practitioner toolkits

6.2 Licensing Model

The Meridian Method™ is available under three licensing structures:

License Type	Rights Granted	Indicative Fee
Practitioner License	Individual use of Meridian ceremonies and job aids within a single organization. Non-transferable. No right to train others or sublicense.	**Included in Allied Skills Program enrollment**
Organizational License	Organization-wide use of the Meridian Method™ across all projects. Includes train-the-trainer rights for internal facilitators. No right to sublicense to clients.	**USD $25,000–$60,000/year depending on org size**
Delivery Partner License	Right to deliver Meridian-certified engagements to client organizations. Includes co-branding rights, certified facilitator program, and access to full IP library.	**USD $40,000–$80,000/year + rev share**

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.

6.3 Certification

LegacyForward will issue the Meridian Method™ Certified Practitioner credential to individuals who complete the Allied Skills training program and pass the Meridian Method™ Practitioner Assessment. The credential is role-specific — a BA earns a different certification path than a QA or PM — and requires annual renewal to remain current with methodology updates.

Organizational and Delivery Partner licensees may nominate internal Meridian Facilitators, who complete an extended certification program enabling them to run Meridian ceremonies and train practitioners within their licensed scope.

6.4 What LegacyForward Does Not Own

LegacyForward does not claim ownership of: the general concept of eval-driven development; the use of confidence thresholds in ML system design; human-in-the-loop development practices as a general principle; or any LLM evaluation framework (RAGAS, LLM-as-judge, etc.) used within Meridian ceremonies. The IP claim is the specific methodology structure, ceremony design, and practitioner framework — not the underlying concepts it organizes.

The Meridian Method™

A LegacyForward IP Publication

legacyforward.ai | © 2025 LegacyForward. All rights reserved.

© LegacyForward. All rights reserved. Meridian Method™ is a trademark of LegacyForward.